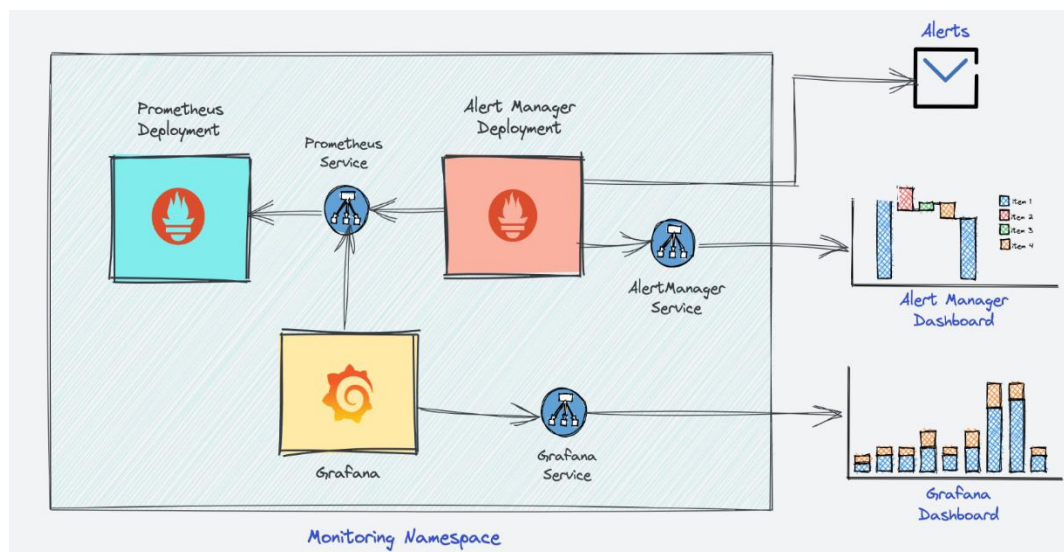


Setup Prometheus Monitoring On Kubernetes Cluster

Prometheus is a high-scalable open-source monitoring framework. It provides out-of-the-box monitoring capabilities for the Kubernetes container orchestration platform. Also, In the observability space, it is gaining huge popularity as it helps with metrics and alerts.

Prometheus Architecture: - The Kubernetes Prometheus monitoring stack has the following components.

1. Prometheus Server
2. Alert Manager
3. Grafana



Step 1: clone the git for Prometheus Kubernetes Manifest Files

```
$ git clone https://github.com/techiescamp/kubernetes-prometheus
```

Step 2: Create a Namespace & Cluster Role

```
$ kubectl create namespace monitoring
```

Step 3: Create a file named clusterRole.yaml and copy the following RBAC role.

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: prometheus
rules:
- apiGroups: [""]
  resources:
    - nodes
    - nodes/proxy
    - services
    - endpoints
    - pods
  verbs: ["get", "list", "watch"]
- apiGroups:
  - extensions
  resources:
  - ingresses
  verbs: ["get", "list", "watch"]
- nonResourceURLs: ["/metrics"]
  verbs: ["get"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: prometheus
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
```

```
  name: prometheus
subjects:
- kind: ServiceAccount
  name: default
  namespace: monitoring
```

Step 4: Create the role using the following command.

```
$ kubectl create -f clusterRole.yaml
```

Step 5: create a config map to externalize Prometheus configurations

All configurations for Prometheus are part of `prometheus.yaml` file and all the alert rules for Alert manager are configured in `prometheus.rules`.

`prometheus.yaml`: This is the main Prometheus configuration which holds all the scrape configs, service discovery details, storage locations, data retention configs, etc)

`prometheus.rules`: This file contains all the Prometheus alerting rules

Step 6: Create a file called `config-map.yaml`

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-server-conf
  labels:
    name: prometheus-server-conf
    namespace: monitoring
data:
  prometheus.rules: |-
    groups:
    - name: devopscube demo alert
      rules:
      - alert: High Pod Memory
        expr: sum(container_memory_usage_bytes) > 1
```

```
    for: 1m
    labels:
      severity: slack
    annotations:
      summary: High Memory Usage
prometheus.yml: |-
  global:
    scrape_interval: 5s
    evaluation_interval: 5s
  rule_files:
    - /etc/prometheus/prometheus.rules
  alerting:
    alertmanagers:
      - scheme: http
        static_configs:
          - targets:
              - "alertmanager.monitoring.svc:9093"
  scrape_configs:
    - job_name: 'node-exporter'
      kubernetes_sd_configs:
        - role: endpoints
      relabel_configs:
        - source_labels: [__meta_kubernetes_endpoints_name]
          regex: 'node-exporter'
          action: keep
    - job_name: 'kubernetes-apiservers'
      kubernetes_sd_configs:
        - role: endpoints
      scheme: https
      tls_config:
        ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
        bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
      relabel_configs:
        - source_labels: [__meta_kubernetes_namespace,
__meta_kubernetes_service_name, __meta_kubernetes_endpoint_port_name]
          action: keep
          regex: default;kubernetes;https
    - job_name: 'kubernetes-nodes'
      scheme: https
      tls_config:
        ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
        bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
      kubernetes_sd_configs:
        - role: node
```

```
relabel_configs:
- action: labelmap
  regex: __meta_kubernetes_node_label_(.+)
- target_label: __address__
  replacement: kubernetes.default.svc:443
- source_labels: [__meta_kubernetes_node_name]
  regex: (.+)
  target_label: __metrics_path__
  replacement: /api/v1/nodes/${1}/proxy/metrics
- job_name: 'kubernetes-pods'
  kubernetes_sd_configs:
  - role: pod
    relabel_configs:
    - source_labels: [__meta_kubernetes_pod_annotation_prometheus_io_scrape]
      action: keep
      regex: true
    - source_labels: [__meta_kubernetes_pod_annotation_prometheus_io_path]
      action: replace
      target_label: __metrics_path__
      regex: (.+)
    - source_labels: [__address__,
__meta_kubernetes_pod_annotation_prometheus_io_port]
      action: replace
      regex: ([^:]+)(?::\d+)?;(\d+)
      replacement: $1:$2
      target_label: __address__
    - action: labelmap
      regex: __meta_kubernetes_pod_label_(.+)
    - source_labels: [__meta_kubernetes_namespace]
      action: replace
      target_label: kubernetes_namespace
    - source_labels: [__meta_kubernetes_pod_name]
      action: replace
      target_label: kubernetes_pod_name
- job_name: 'kube-state-metrics'
  static_configs:
  - targets: ['kube-state-metrics.kube-system.svc.cluster.local:8080']
- job_name: 'kubernetes-cadvisor'
  scheme: https
  tls_config:
    ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
    bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
  kubernetes_sd_configs:
  - role: node
```

```

    relabel_configs:
      - action: labelmap
        regex: __meta_kubernetes_node_label_(.+)
      - target_label: __address__
        replacement: kubernetes.default.svc:443
      - source_labels: [__meta_kubernetes_node_name]
        regex: (.+)
        target_label: __metrics_path__
        replacement: /api/v1/nodes/${1}/proxy/metrics/cadvisor
    - job_name: 'kubernetes-service-endpoints'
      kubernetes_sd_configs:
        - role: endpoints
      relabel_configs:
        - source_labels:
[__meta_kubernetes_service_annotation_prometheus_io_scrape]
          action: keep
          regex: true
        - source_labels:
[__meta_kubernetes_service_annotation_prometheus_io_scheme]
          action: replace
          target_label: __scheme__
          regex: (https?)
        - source_labels:
[__meta_kubernetes_service_annotation_prometheus_io_path]
          action: replace
          target_label: __metrics_path__
          regex: (.+)
        - source_labels: [__address__,
__meta_kubernetes_service_annotation_prometheus_io_port]
          action: replace
          target_label: __address__
          regex: ([^:]+)(?::\d+)?;(\d+)
          replacement: $1:$2
        - action: labelmap
          regex: __meta_kubernetes_service_label_(.+)
        - source_labels: [__meta_kubernetes_namespace]
          action: replace
          target_label: kubernetes_namespace
        - source_labels: [__meta_kubernetes_service_name]
          action: replace
          target_label: kubernetes_name

```

Step 7: Execute the following command to create the config map in Kubernetes.

```
$ kubectl create -f config-map.yaml
```

Step 8: Check wheather it is created or not

```
$ kubectl describe cm prometheus-server-conf -n monitoring
```

```
~$ kubectl describe cm prometheus-server-conf -n monitoring
Name:      prometheus-server-conf
Namespace: monitoring
Labels:    name=prometheus-server-conf
Annotations: <none>

Data
====
prometheus.rules:
----
groups:
- name: devopscube demo alert
  rules:
  - alert: High Pod Memory
    expr: sum(container_memory_usage_bytes) > 1
    for: 1m
    labels:
      severity: slack
    annotations:
      summary: High Memory Usage
prometheus.yml:
----
global:
  scrape_interval: 5s
  evaluation_interval: 5s
rule_files:
- /etc/prometheus/prometheus.rules
alerting:
  alertmanagers:
  - scheme: http
    static_configs:
    - targets:
      - "alertmanager.monitoring.svc:9093"
scrape_configs:
- job_name: 'node-exporter'
  kubernetes_sd_configs:
  - role: endpoints
  relabel_configs:
  - source_labels: [ __meta_kubernetes_endpoints_name ]
    regex: 'node-exporter'
    action: keep
- job_name: 'kubernetes-apiservers'
  kubernetes_sd_configs:
  - role: endpoints
  scheme: https
```

Step 9: Create a file named **prometheus-deployment.yaml** and copy the following contents onto the file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: prometheus-deployment
```

```
namespace: monitoring
labels:
  app: prometheus-server
spec:
  replicas: 1
  selector:
    matchLabels:
      app: prometheus-server
  template:
    metadata:
      labels:
        app: prometheus-server
    spec:
      containers:
        - name: prometheus
          image: prom/prometheus
          args:
            - "--storage.tsdb.retention.time=12h"
            - "--config.file=/etc/prometheus/prometheus.yml"
            - "--storage.tsdb.path=/prometheus/"
          ports:
            - containerPort: 9090
          resources:
            requests:
              cpu: 500m
              memory: 500M
            limits:
              cpu: 1
              memory: 1Gi
          volumeMounts:
            - name: prometheus-config-volume
              mountPath: /etc/prometheus/
            - name: prometheus-storage-volume
              mountPath: /prometheus/
      volumes:
        - name: prometheus-config-volume
          configMap:
            defaultMode: 420
            name: prometheus-server-conf

        - name: prometheus-storage-volume
          emptyDir: {}
```


Step 10: Create a deployment on monitoring namespace using the above file.

```
$ kubectl create -f prometheus-deployment.yaml
```

Step 11: You can check the created deployment using the following command.

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
Prometheus-deployment	1/1	1	1	5d20h

Step 12: Create a file named `prometheus-service.yaml` and copy the following contents. We will expose Prometheus on all kubernetes node IP's on port 30000.

```
apiVersion: v1
kind: Service
metadata:
  name: prometheus-service
  namespace: monitoring
  annotations:
    prometheus.io/scrape: 'true'
    prometheus.io/port: '9090'
spec:
  selector:
    app: prometheus-server
  type: NodePort
  ports:
    - port: 8080
      targetPort: 9090
      nodePort: 30000
```

Step 13: Create the service using the following command.

```
$ kubectl create -f prometheus-service.yaml --namespace=monitoring
```

Step 14: Now we can access to use IP address and Nodeport check in the Browsers

Step 15: **kube state metrics setup** and Clone the Github repo

```
$ git clone https://github.com/devopscube/kube-state-metrics-configs.git
```

Step 16: Create all the objects by pointing to the cloned directory.

```
$ kubectl apply -f kube-state-metrics-configs/
```

Step 17: Check the deployment status using the following command.

```
$ kubectl get deployments kube-state-metrics -n kube-system
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
kube-state-metrics	1/1	1	1	5d4h

Step 18: Setup Node Exporter on Kubernetes Create a file name **daemonset.yaml** and copy the following content.

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  labels:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: node-exporter
```

```
name: node-exporter
namespace: monitoring
spec:
  selector:
    matchLabels:
      app.kubernetes.io/component: exporter
      app.kubernetes.io/name: node-exporter
  template:
    metadata:
      labels:
        app.kubernetes.io/component: exporter
        app.kubernetes.io/name: node-exporter
    spec:
      containers:
        - args:
            - --path.sysfs=/host/sys
            - --path.rootfs=/host/root
            - --no-collector.wifi
            - --no-collector.hwmon
            - --collector.filesystem.ignored-mount-
points=^/(dev|proc|sys|var/lib/docker/.+|var/lib/kubelet/pods/.+)($/|/)
            - --collector.netclass.ignored-devices=^(veth.*)$
          name: node-exporter
          image: prom/node-exporter
          ports:
            - containerPort: 9100
              protocol: TCP
          resources:
            limits:
              cpu: 250m
              memory: 180Mi
            requests:
              cpu: 102m
              memory: 180Mi
          volumeMounts:
            - mountPath: /host/sys
              mountPropagation: HostToContainer
              name: sys
              readOnly: true
            - mountPath: /host/root
              mountPropagation: HostToContainer
              name: root
              readOnly: true
          volumes:
```

```
- hostPath:
  path: /sys
  name: sys
- hostPath:
  path: /
  name: root
```

Step 19: Deploy the daemon set using the kubectl command.

```
$ kubectl create -f daemonset.yaml
```

Step 20: List the daemon set in the monitoring namespace and make sure it is in the available state.

```
$ kubectl get daemonset -n monitoring
```

Step 21: Create a file named `service.yaml` and copy the following contents.

```
kind: Service
apiVersion: v1
metadata:
  name: node-exporter
  namespace: monitoring
  annotations:
    prometheus.io/scrape: 'true'
    prometheus.io/port: '9100'
spec:
  selector:
    app.kubernetes.io/component: exporter
    app.kubernetes.io/name: node-exporter
  ports:
    - name: node-exporter
      protocol: TCP
      port: 9100
      targetPort: 9100
```

Step 22: Create the service.

```
$ kubectl create -f service.yaml
```

Step 23: Now, check the service's endpoints and see if it is pointing to all the daemon set pods.

```
$ kubectl get endpoints -n monitoring
```

```
→ ~ kubectl get daemonset -n monitoring
NAME          DESIRED  CURRENT  READY  UP-TO-DATE  AVAILABLE  NODE SELECTOR
node-exporter  3        3        3      3           3          <none>
→ ~ kubectl get service -n monitoring
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)          AGE
grafana       NodePort    10.245.13.187 <none>       3000:32000/TCP   15h
node-exporter ClusterIP    10.245.72.166 <none>       9100/TCP         9h
prometheus-service NodePort    10.245.2.190  <none>       8080:30000/TCP   15h
→ ~ kubectl get endpoints -n monitoring
NAME          ENDPOINTS                                          AGE
grafana       10.244.0.160:3000                                15h
node-exporter 10.244.0.223:9100,10.244.1.183:9100,10.244.1.69:9100 9h
prometheus-service 10.244.0.195:9090                                15h
→ ~
```

As you can see from the above output, the node-exporter service has three endpoints. Meaning three node-exporter pods running on three nodes as part of Daemon set. Grafana

Grafana Setup on standalone server

Step 1: Install Required Packages

```
$ sudo apt-get install -y apt-transport-https software-properties-common
wget
```

Step 2: Create Directory for GPG Key

Create a directory to store the GPG key used to authenticate the Grafana repository.

```
$ sudo mkdir -p /etc/apt/keyrings/
```

Step 3: Download and Store Grafana GPG Key

```
$ wget -q -O - https://apt.grafana.com/gpg.key | gpg --dearmor | sudo tee  
/etc/apt/keyrings/grafana.gpg > /dev/null
```

Step 4: Add Grafana Stable Repository

Add the stable Grafana repository to your system's sources list.

```
$ echo "deb [signed-by=/etc/apt/keyrings/grafana.gpg]  
https://apt.grafana.com stable main" | sudo tee -a  
/etc/apt/sources.list.d/grafana.list
```

Step 5: (Optional) Add Grafana Beta Repository

```
$ echo "deb [signed-by=/etc/apt/keyrings/grafana.gpg]  
https://apt.grafana.com beta main" | sudo tee -a  
/etc/apt/sources.list.d/grafana.list
```

Step 6: Verify the Sources List

Navigate to the sources list directory and review the newly added Grafana repositories.

```
$ cd /etc/apt/sources.list.d/
```

```
sudo vim grafana.list
```

Step 7: Update Package Lists

```
$ sudo apt-get update
```

Step 8: Install Grafana

Now, install the Grafana package.

```
$ sudo apt-get install grafana
```

Step 10: Start Grafana Server

Start the Grafana server service.

```
$ sudo systemctl start grafana-server.service
```

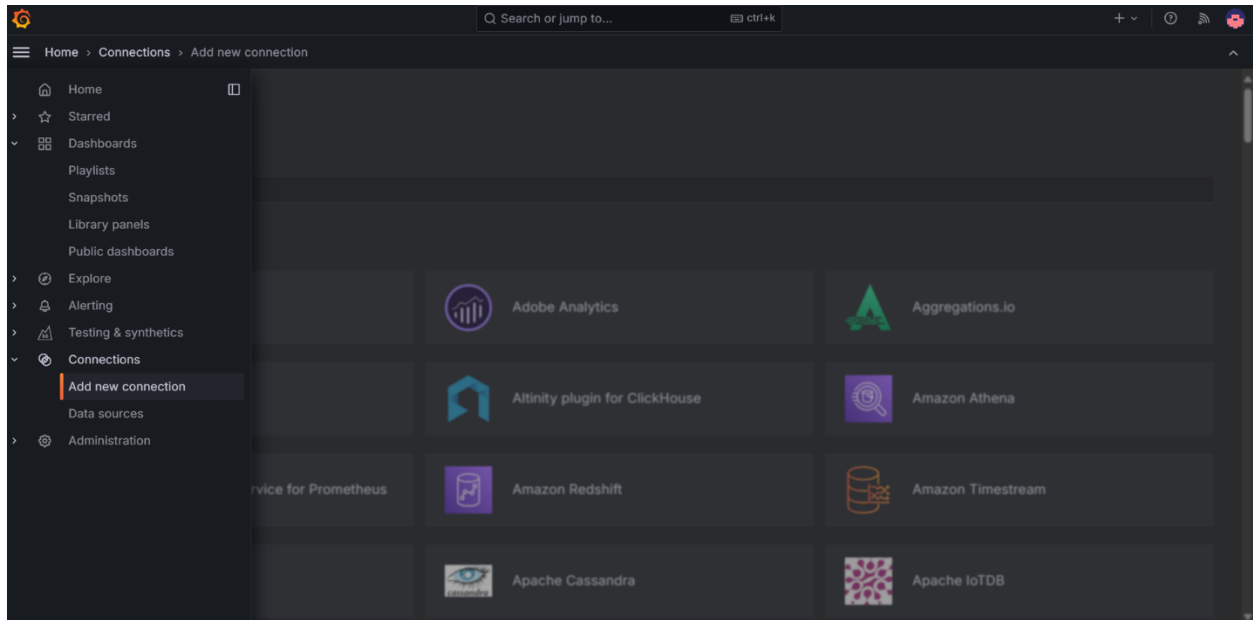
Step 11: Check Grafana Server Status

Check the status of the Grafana server to ensure it's running correctly.

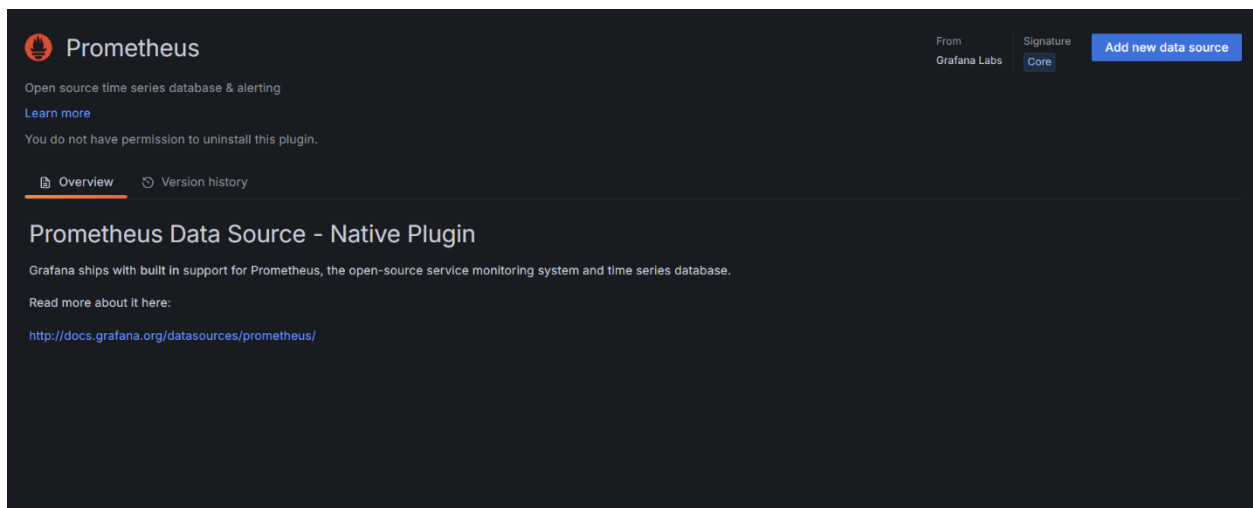
```
$ sudo systemctl status grafana-server.service
```

Step 12 : login with user and password

Step 13: add connection search the Prometheus and it



Step 14: add data source in Prometheus



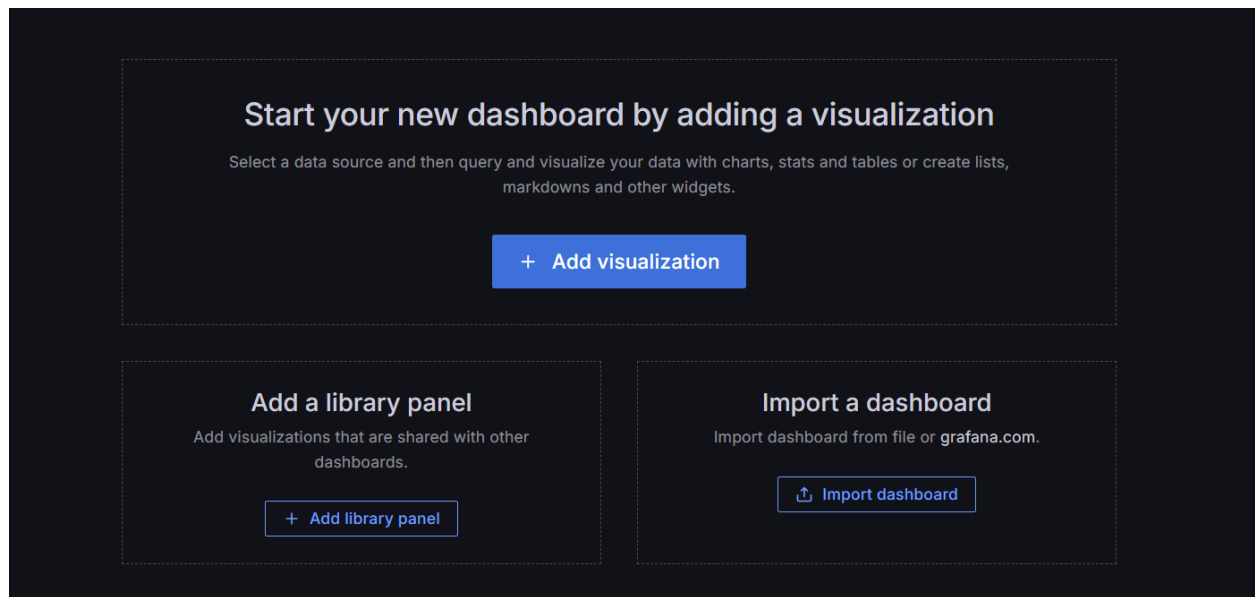
Enter the fill details regarding Prometheus

Connection

Prometheus server URL *

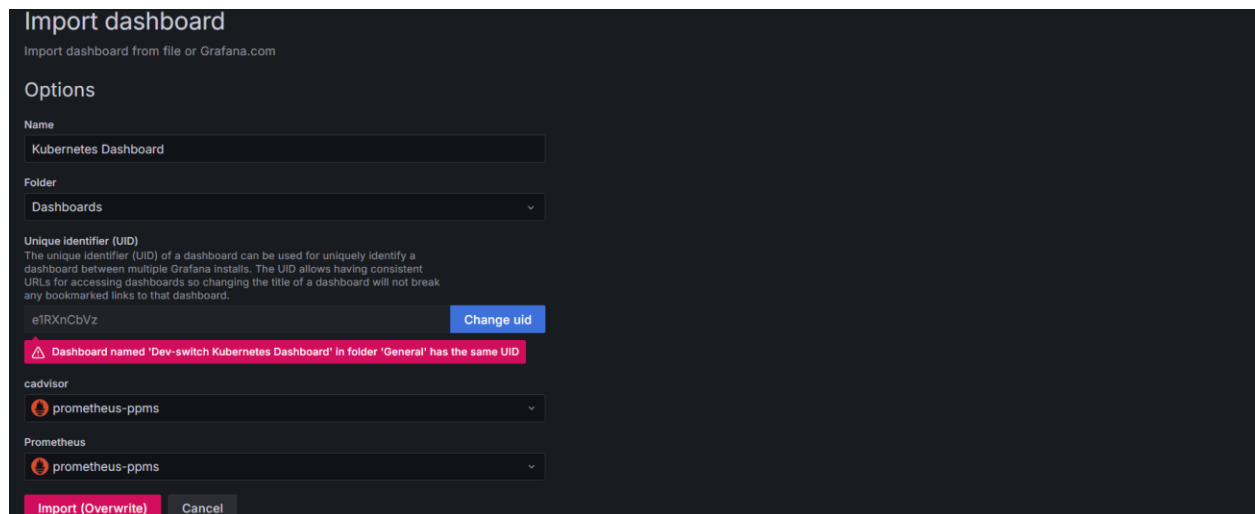
Step 15: Goto dashboard add now dashboard

Step 16: go to import dashboard

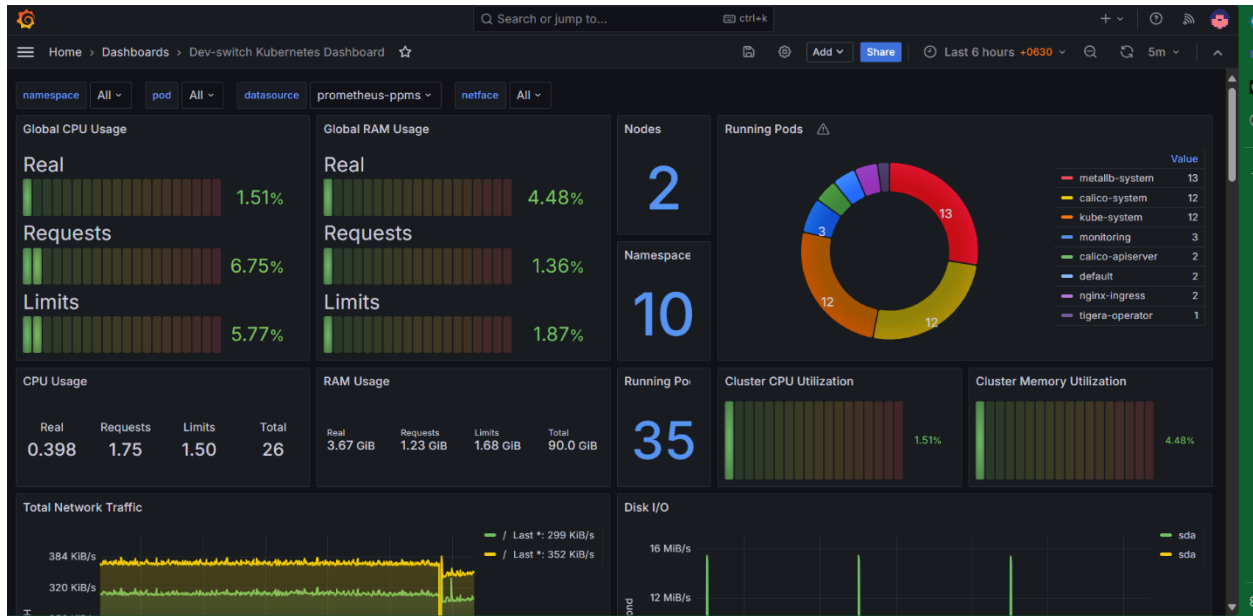


Step 17: In google search Grafana dashboard and download the dashboard as per requirement

Step 18: import the dashboard from download give Prometheus connection in dashboard



Step 19:- we can see grafana dashboard



<https://devopscube.com/setup-prometheus-monitoring-on-kubernetes/>

<https://devopscube.com/setup-kube-state-metrics/>

[How To Setup Prometheus Node Exporter On Kubernetes \(devopscube.com\)](#)